


From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · PROJECT BRIEF

12-Week Hybrid Agentic System Pilot

Project Brief



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · PROJECT BRIEF


1. Executive Summary

Finexa's product function is strained by feedback volume arriving across six disconnected channels (CS p. 1). The cost is not only slow synthesis: backlog items are not traceable to verified customer needs, and prioritization runs on anecdote rather than evidence (CS p. 2).

This brief recommends a **12-week pilot in the Payments domain — a Hybrid Agentic System. It combines an autonomous LLM reasoning core with Appsmith UI and workflow automation, deterministic infrastructure, and MCP-driven analytics.**

The Process diagram shows each step and each table in the repository created/updated.

Picture 1: Process Diagram Click on the diagram to open the full view.



Key Architecture Capabilities:

1) Pure Cognitive Focus (Decoupled Infrastructure)

- **Upstream DLP (Non-AI):** De-identification and PII scrubbing run deterministically before text touches the LLM.
- **Core LLM Engine:** Focuses purely on cognitive tasks — atomic 1-to-N feedback deconstruction, ENUM schema extraction, structuring and summarization of feedback, vector search and shortlisting of recurring issues, and structured Jira payload generation in BigQuery staging.
- **Downstream Sync (Non-AI):** Once a human approves Gate 3, a BigQuery wrapper script automatically dispatches approved tickets to Jira.

2) Native Tool & Operational Integration (Appsmith + MCP)

- **Appsmith Operational Hub:** Replaces manual spreadsheets with a unified UI featuring workflow automation and human-in-the-loop Gate 1–3 controls.
- **Vector RAG Matching:** Directly queries BigQuery/Vector databases to surface recurring issues automatically and reads and compares them with the insight.
- **/ask_data MCP Agent:** Connects to BigQuery via Model Context Protocol, allowing to run natural-language operational queries, track override rates, and evaluate product metrics instantly without writing manual SQL or exporting CSVs.

3) Human-in-the-Loop Governance (Gate 1–3 Architecture)

Instead of running as an unmonitored "black box," the system acts as an intelligent copilot that pre-builds structured artifacts for human validation:

- **Gate 1:** Review/approve atomic problem splits, summary, ENUM tags (severity, feature area etc.). Review/approve RAG candidate matches.
- **Gate 2:** Review/approve scoring, approve opportunities for drafting.
- **Gate 3:** Review/approve pre-drafted Jira payloads and estimated complexity.

2. Business Need

Over the last two quarters, Finexa has seen rising customer complaints across payment and dispute workflows alike (CS p. 3) — merchants who don't know whether a payment is pending, failed, or settled, or what a dispute actually requires of them.

The organization already has strong customer access: 25 interviews, 100+ support tickets, ops escalations, strategic account requests, and Jira issues about recurring confusion (CS p. 3). What it lacks is a way to turn that volume into a shared, current, evidence-backed understanding — today, insights live in scattered notes and spreadsheets, prioritization runs on anecdote, and backlog items regularly arrive at engineering incomplete (CS p. 4).

3. Objective & Requirements

Objective: move the Merchant Experience Product Group from fragmented feedback to evidence-based delivery planning (CS p. 4), without adding headcount or introducing unmanaged tooling.

Three conditions, drawn from the stakeholder guardrails below, bound how this is achieved: strict **security and compliance** control, **no unreviewed AI output** reaching the engineering backlog, and **no removal of the PM's decision authority**.

Key Stakeholder Requirements & Guardrails:

- **Elena Vance (Head of Product):** demands fast turnaround and a clear, auditable evidence trail during roadmap reviews (CS p. 6).
- **Marcus Bauer (Senior PM):** wants relief from manual synthesis, but insists the AI stay an assistant, never his replacement (CS p. 7).
- **Priya Patel (Engineering Lead):** will resist adoption unless strict, non-negotiable review checkpoints keep incomplete or hallucinated tickets out of her backlog (CS p. 7).
- **Simone Dubois (Risk & Compliance):** will block the pilot outright without strict security controls, GDPR compliance, and full auditability (CS p. 8).
- **Liam O'Connor (Support Ops):** needs assurance that operational escalations won't be quietly buried under feature requests (CS p. 8).

Key Constraints:

The wider team has already identified several existing constraints any pilot must work within (CS p. 9):

- **Regulated data** — the work involves regulated financial data, handled within strict controls.
- **A short runway** — limited time for implementation.
- **Uneven data quality** — source material varies considerably across interviews, tickets, and notes.
- **Human oversight** — the process depends on human validation at key points.
- **Adoption** — uptake across teams cannot be assumed.

4. Assumptions

Platform & Technical

- Finexa's existing AI platform can host custom reusable commands, and BigQuery is already Finexa's own data platform — not something introduced *for* this pilot.
- Finexa's Jira workspace uses Fix Versions with real release dates. This is what lets a ticket marked "Done" be told apart from one actually shipped — without it, regression detection has no reliable signal to key off.
- Finexa's Jira supports API-based ticket creation via a service account, not manual entry only.

Data sources & Legal

- The six feedback channels are exportable into BigQuery in reasonably clean form, not, say, unstructured audio with no transcription step already in place.
- Merchant account tier (Enterprise/SMB) already exists as a queryable field on Finexa's merchant records.
- A merchant's `merchant_id` is non-personal data for corporate accounts; for sole proprietors it may identify an individual and needs Legitimate Interest sign-off (TBD with Simone).


Operational & Adoption

- Marcus and Priya have real bandwidth to review consistently across the full 12 weeks — the pilot's evidence is only as good as review actually happening.
- Appsmith access is provisioned for both reviewers.

5. Core Workflow Description.

5.1 The Process diagram shows each step and each table in the repository created/updated with four sequential commands — **/process-feedback**, **/log-feedback**, **/prioritize-opportunities**, **/draft-user-story** — plus **/ask-data** running in parallel for read-only queries, all inside Finexa's existing **AI platform**, writing to and reading from **BigQuery** as the SSOT and using **Appsmith** as the workflow orchestrator and UI.

Picture 2: Process Diagram. Click on the diagram to open the full view.



1. **Google DLP (Data Loss Protection)** — scrubbing PII and sensitive financial information downstream prevents payment/merchant data reaching the model.
2. **/process-feedback (Assist)** — splits one transcript or ticket into distinct problems, each with a summary, pain points, root cause, and a provisional theme. Thin material halts and asks for clarification rather than guessing.
3. **/log-feedback (Assist)** — proposes whether this is the same issue as something already confirmed, or something new.
4. **Gate 1 (mandatory at this stage)** — confirms the reading and the grouping. A *verification* task: is this the same theme as something already validated? Because that has a checkable correct answer, a sustained agreement rate is real evidence Gate 1 could eventually move to spot-checks — the only one of the three gates designed to.
5. **/prioritize-opportunities (Augment)** — reads a priority score already computed by a separate, deterministic formula and explains it in plain language. It never calculates the number itself.
6. **Score Opportunity** — five inputs, each pulled from evidence already sitting in the system, no separate manual step. **Customer value** weighs how many Enterprise vs. SMB merchants are affected. **Urgency** combines the single worst severity reported — one "Critical" call from a customer counts, it isn't averaged down by milder reports — with how often the issue recurs. **Strategic alignment** and **operational risk** are read from configuration set ahead of time by Elena's team and engineering, not decided fresh for each item. **Compliance sensitivity** works the same way. Each input carries a tunable weight, and the score is their weighted sum — arithmetic, not a model call. **Implementation complexity** is deliberately not one of the five; it's assigned manually at Gate 3 instead, once a real draft exists to judge, because no reliable signal exists to estimate it at this stage.

7. **Gate 2 (permanent)** — confirms this specific scoring opportunity is worth spending drafting effort on now. A *judgment* task, not a verification one: by this point the item is already a validated, scored opportunity, and deciding what matters right now is shaped by capacity and timing that shift week to week — there's no single correct answer an agreement rate could ever converge toward, so this gate never graduates.
8. **/draft-user-story (Assist)** — converts an approved opportunity into a problem statement, user story, acceptance criteria, and evidence-linked rationale.
9. **Gate 3 (permanent)** — confirms the draft against the original evidence and the suggested implementation complexity.
10. **The deterministic Wrapper Code** — the only component holding Jira credentials — calls Jira's API and posts selected Jira Drafts to Jira.
11. **/ask-data (Augment)** — plain-language, read-only querying of stored insights, running in parallel to the sequence above, any time.

5.2 Capabilities of the Architecture:

- **PII is unreachable by design.** Cloud DLP de-identifies feedback at ingestion, and the AI's service account holds no credential that could retrieve raw source data — it is not instructed to avoid PII, it has no path to it. The residual risk is free-text residue that DLP misses, controlled by pattern validation before ingestion and ongoing sampling.
- **Parsing and structuring with an LLM into a consistent insight schema** and writing to BigQuery, gives both queryability and a foundation for vector search.
- **Vector search** for similar issues + linked opportunities is the highest-leverage part. The same complaint arriving in an interview, a support ticket and a CRM note becomes **one prioritized opportunity carrying three pieces of evidence**, rather than three unconnected backlog entries.
- **Three mandatory human gates sit inside the sequence.** The AI drafts and proposes; a product manager or the engineering lead decides. No command holds Jira credentials, so no AI output can reach engineering unreviewed — a structural guarantee rather than a policy commitment.
- **The audit trail is inherited, not built.** Every gate decision is a stored row and BigQuery logs data access by default — so auditability is a property of the platform Finexa already runs, not custom logging code that must be maintained and trusted to have been switched on.
- **Priority is a formula, not a judgment.** The score is arithmetic over five of the six criteria the case names; implementation complexity is deliberately excluded — it's assigned manually at Gate 3 once a real draft exists, not estimated by the formula. Only the score's explanation is AI-written. Because the weights are parameters, a systematic disagreement at review identifies which weight is wrong and who can change it. An AI-generated score could be disputed, but not tuned.

- **Dedicated UI layer.** People need to see context side-by-side, compare similar opportunities, edit structured fields, and act with one or two clicks.
- **AI Auditability (gate_decisions):** Storing `ai_proposed_value` alongside `pm_final_value` gives you an immediate dataset to measure prompt drift, model accuracy, and human override rates over time.
- **Regression Intelligence (regression_flag + last_insight_at):** Capturing when incoming feedback maps to a resolved or existing canonical opportunity enables automatic detection of post-release regressions or incomplete fixes.

5.3 BigQuery Schema - *Click to open the full view.*

5.4 User Interface

All three gates are executed through a single Appsmith interface rather than across separate tools. Each screen presents the item under review alongside the evidence behind it, with AI-proposed values as editable fields. The three views below show Gate 1, Gate 2, and Gate 3 in sequence.

Human Gate 1:

Picture 4: Insight linked to the source ticket and opportunities. Click on the diagram to open the full view.

The screenshot displays the 'Finexa · Feedback Intelligence' interface. The main content area is titled 'Insight #INS-0234' with tags 'Support ticket' and 'Enterprise', dated '14 Mar 2025'. The status is 'Pending Review'.

Summary: Merchant cannot see why a payment failed, only a generic error message shown in the dashboard.

Problem details:

- Understand why each payment failed
- Recover in the dashboard

Jobs & goals:

- No failure reason or code
- Know what action is in the UI
- Support what open internal tools to investigate

Pain points: Failure codes are logged, but not exposed in the merchant-facing payment overview

Classification:

- Workflow: **Payment status and failure reasons**
- Product area: **payments_status**
- Severity: **High** (blue dot), **Elevated** (blue dot)

Source ticket · ST-49281

Ticket ID: **ST-49281** Created: 12 Mar 2025 · 14:22

Customer: **Acme Payments GmbH (EU · Enterprise)** Channel: **Email → Support**

Original message (de-identified)

"When a payment fails we only see 'Payment failed'. We need the actual decline reason (insufficient funds, card expired, 3DS failure, etc.) so our team can tell the merchant what to do next. Right now we have to open internal tools for every case."

Priority: **High** Status: **Open**

Matching opportunities

- OP-011 · Payment failures lack clear reasons** 0.87
9 linked insights · High priority · for drafting
[View full →](#)
- OP-019 · Generic error messages in checkout** 0.72
4 linked insights · Medium priority · monitoring
[View full →](#)

Opportunity OP-011 · Payment failures lack clear reasons in merchant UI

Workflow: **Payment status** Feature area: **payments_status** **for drafting**

Evidence summary

0.82 High priority

High impact: 9 reports including 4 enterprise. Recent escalations and high operational risk if unsolved.

Priority

9 linked insights

Severity: 4 High · 3 Medium · 2 Low
Segments: 4 Enterprise · 5 SMB

Recent insights

Insight	Segment	Severity	Excerpt
INS-0210	Enterprise	High	Cannot see why a payme...
INS-0187	Enterprise	High	Need explicit failure reas...

Linked Jira tickets (2)

- PAY-421 · Expose failure reason in payment details** **In progress**
Created 03 Feb 2025 · Assignee: Anna K. 5 pts
Add decline code + human-readable reason to the merchant payment detail view. Include next-best-action suggestion.
payments merchant-ui
- PAY-397 · Improve error messaging for failed payments** **Resolved**
Created 18 Nov 2024 · Assignee: Tomas R. 3 pts
Replaced generic "Payment failed" with more specific messages for the top 5 decline reasons.
checkout shipped 2.3.0

Warning: This opportunity has new insights after a previous fix. Treat as potential regression.

Human Gate 2:

Picture 5: Opportunity is scored and explained. PM selects opportunities for drafting. Click on the diagram to open the full view.

Finexa · Feedback Intelligence

Gate 2 · Marcus

Prioritized Opportunities

Ranked by score · highest first

#1 OP-011

Regression

Approved

0.87

score

Payment failures lack clear reasons in merchant UI

9 insights · 4 Enterprise · 5 SMB

#2 OP-019

0.74

score

Generic error messages in checkout flow

6 insights · 2 Enterprise · 4 SMB

#3 OP-024

0.61

score

Merchant dashboard lacks payment method health indicators

4 insights · 1 Enterprise · 3 SMB

#4 OP-031

0.53

score

Refund status not visible to merchants in real time

3 insights · 0 Enterprise · 3 SMB

OP-019 · prioritized

Generic error messages in checkout flow

0.74

priority score

Linked insights

6

Severity mix

2H · 3M · 1L

Segments

2 Enterprise · 4 SMB

SCORE EXPLANATION

Recurring feedback from both SMB and enterprise. Lower enterprise weight than OP-011 but still material. Several tickets already linked in previous quarters.

Last scored 5 hours ago

Theme

Checkout error clarity

Workflow

Checkout

Feature area

checkout_flow

Approve for drafting

Skip for now

Human Gate 3:

Picture 6: Jira Draft is checked, complexity is estimated, the draft can be confirmed for Jira. Click on the diagram to open the full view.

Finexa · Feedback Intelligence

Gate 3 · Review Drafts

Jira Drafts

Approved opportunities · ready for review

#1 OP-011

Regression

Confirmed

5 pts

complexity

Expose failure reason in payment details

9 insights · score 0.87

#2 OP-019

3 pts

complexity

Improve error messaging in checkout

6 insights · score 0.74

#3 OP-024

8 pts

complexity

Add payment method health indicators to dashboard

4 insights · score 0.61

LINKED OPPORTUNITY

OP-019

Generic error messages in checkout flow

0.74

score

Checkout error clarity

Checkout

checkout_flow

6 insights

2 Enterprise · 4 SMB

Jira Draft · DRAFT-002

pending_draft

TITLE

Improve error messaging in checkout

PROBLEM STATEMENT

During checkout, customers and merchants see only generic error messages when a payment method is declined. This leads to abandoned checkouts and unnecessary support contacts.

USER STORY

As a customer (or merchant assisting a customer), I want clear, actionable error messages at checkout so I know how to fix the payment problem and complete the purchase.

ACCEPTANCE CRITERIA

Replace generic "Payment failed" with specific messages for the top 5 decline reasons

Messages include a short recovery action (e.g. "Try another card" or "Contact your bank")

Works on both desktop and mobile checkout

COMPLEXITY (STORY POINTS)

1

2

3

5

8

13

AI suggested: 3

Confirm draft & complexity

Edit draft

Send back

12

5.5 /ask_data:

With this volume of clean and structured data continuously flowing through the pipeline, Finexa can dynamically execute three distinct layers of operational analytics: Product Analytics to trace macro feature requirements, Process Analytics to monitor team workflow optimization, and AI Operations Analytics to audit core model efficacy.

To deliver these insights natively, the platform couples an interactive Appsmith frontend interface with a conversational ask_data Analytical Agent powered by the open-standard Model Context Protocol (MCP). This dual architecture allows the organization to seamlessly deploy both persistent, real-time analytics dashboards for standard executive monitoring and ad-hoc, natural-language reporting for immediate, flexible data exploration.

Picture 7: Click [/ASK DATA](#) to open the full report prototype.



6. Scope of the Pilot

In-Scope (Phase V1 — 12 Weeks)

- **Project Documentation & Governance** — produce Project Brief, Stakeholder Map + RACI, Risk Register, Monitoring Plan. Create/update an internal AI application register tracking.
- **Technical Documentation:** define requirements, produce architecture/schema/command documentation
- **Baseline & measurement:** baseline data report + PM/Eng/Support/stakeholder measures + targets
- **Five AI commands:** confirm requirements/guardrails, configure/test, assess quality + adoption
- **BigQuery schema:** confirm data/governance requirements, create tables/relationships/status fields
- **Config tables:** agree criteria/weights/policy rules, create scoring + taxonomy configs
- **Appsmith UI:** define user journeys/approval roles, build command + review screens
- **Human gates:** agree responsibilities, implement Gate 1/2/3 workflows with decision logs
- **Jira wrapper + sync:** confirm access boundaries, build credential-isolated wrapper, status sync
- **Source data ingestion:** confirm baseline dataset, load + de-identify interviews and tickets
- **Monitoring plan (pilot-level):** define metrics/collection, instrument logs + measurement points. Final project success assessment against the baseline.

Out of scope (V1) — Deferred to enterprise rollout (V2+)

- Enterprise AI governance expansion (standing committees, org-wide AI inventory/register)
- Fine-tuning / continual learning loops (training pipelines, RLHF, automated retraining)
- Enterprise monitoring platform + fully formal 24/7 incident operations
- Autonomous orchestration/triggers (e.g., n8n)
- Full channel coverage / org-wide rollout beyond the pilot domain and dataset
- Executive dashboards / Slack bots beyond what's required for pilot evaluation
- Deep UI productization (full UX research, accessibility/localization hardening)
- Organization-wide change management (company-wide training + adoption KPIs; HR/performance implications beyond pilot-safe logging)

The deferred capabilities outlined in the 'Out of Scope' section serve as the foundation for the Phase V2 Enterprise Rollout, which will be considered upon the successful validation of the Phase V1 pilot metrics against the historical baseline. Once this success threshold is reached and executive approval is secured, the current operational boundaries can expand, shifting all V1 exclusions directly into active, in-scope delivery requirements.

7. Stakeholder MAP and RACI-matrix.


Full details in the Stakeholder MAP and RACI-matrix.

Picture 4: Stakeholder Map and RACI Matrix. Click on the diagram to open the full view.

01 / STAKEHOLDER DIRECTORY		02 / RACI MATRIX		
STAKEHOLDER	POWER	INTEREST	STRATEGY	KEY PRIORITIES & GOVERNANCE FOCUS
Elena Vance HEAD OF PRODUCT	HIGH	HIGH	MANAGE CLOSELY	ROI, overall business value, roadmap alignment, scope approval, and governance risk.
Marcus Bauer SENIOR PM	MEDIUM	HIGH	MANAGE CLOSELY	Product specifications, command accuracy, prompt scope, and target baseline metrics.
Priya Patel ENG LEAD	HIGH	HIGH	MANAGE CLOSELY	System architecture, BigQuery schema, technical debt, UI performance, and integration reliability.
Simone Dubois RISK & COMPLIANCE	HIGH	MEDIUM	KEEP SATISFIED	PII de-identification, Cloud DLP compliance, audit decision logs, and regulatory boundaries.
Liam O'Connor SUPPORT OPS	MEDIUM	HIGH	KEEP INFORMED	Operational usability, source data ingestion accuracy, and end-user support workflow efficiency.
External Consultant IMPLEMENTATION PARTNER	LOW	HIGH	KEEP INFORMED	Clear technical scope, rapid issue unblocking, pipeline execution, and sprint velocity.

8. Timeline & Pilot Structure

Picture 5: Project Implementation Roadmap. Click on the diagram to open the full view.



Weeks 1–4 (Foundation): DLP template sign-off, schema build, baseline capture, stakeholder survey, stakeholder RACI agreement, Appsmith configuration, processing of the existing payments data.

Weeks 5–10 (Pilot Execution): live processing of daily Payments feedback.

Weeks 11–12 (Assessment): measured results against baseline; stakeholder confidence Survey, Proceed / Iterate / Stop decision.

9. Success Criteria & Metrics

Full details in the Monitoring Plan

- A measurable reduction in time spent on manual synthesis.
- Greater consistency in product documentation.
- Improved traceability from a customer insight through to the Jira item it informs.
- A faster turnaround from discovery to backlog readiness.
- Stronger stakeholder confidence in how priorities are set (CS p. 9).

10. Key Risks Summary

These are the risks stakeholders raised directly (CS p. 6–8).

Risk	What could go wrong	Control / safeguard
Backlog flooding / low-quality tickets	Engineering receives many inaccurate/incomplete Jira items	Human approval before “ready for backlog”
Over Reliance on AI	Teams accept drafts without critical thinking	Mandatory evidence citations, human sign-off
Privacy/confidentiality leakage	Direct identifiers or unnecessary financial details enter system	Sanitize before ingestion, store only customer reference
Bias in synthesis/prioritization	Important ops/compliance issues are deprioritized	Explainable scoring that can be tuned + human approval
Adoption stalls	Teams ignore the process, return to spreadsheets	Friendly user interface, clear SOP, fast visible wins

The full analysis and resulting top-5 risks ranking are in the Project's Risks Hub.

Picture 5: Finexa Pilot Executive Risk & Activity Hub. Click to review.

Risk Cockpit

Risk Resolution Activity Hub

Executive Analytics

Taxonomy & Scoring Framework

Top Risks Headlist

ID	Risk Title	Severity
RE-01	Data Privacy & PII Leakage	20
TE-02	Human Gate Rubber-Stamping / Cognitive Fatigue	16
RE-02	Works Council Exposure (BetrVG §87)	15
OP-01	Operational Bandwidth & Gate Latency	15
TE-06	Total Infrastr./ API Outage	15

11. Governance & Compliance

- **Regulated Financial Data:** protected via Cloud DLP de-identification, EU-only data residency, and a defined retention policy on raw files (CS p. 9).
- **EU AI Act Alignment:** classified as an internal drafting/administrative tool, falling outside the high-risk tier — it makes no autonomous decisions affecting an individual's credit, employment, or rights.
- **GDPR Compliance:** relies on Legitimate Interest for feedback ingestion and merchant_id resolution, pending sign-off. merchant_id stays plain, not tokenized, and never travels past the point where the account tier is resolved at Gate 1; every table a PM or the AI can see downstream carries the tier only, never the merchant's identity. Individual names, emails, and phone numbers within free text are what Cloud DLP redacts, before any structured table exists.
- **Works Council (BetrVG §87):** decision logs capture reviewer **role** (e.g., Payment_PM), not individual employee IDs, to avoid performance-monitoring concerns.
- **Data Protection Impact Assessment:** current assessment is none required (pending confirmation). DPIA-relevant triggers concern processing *about natural persons* — scoring, profiling, evaluation. This pipeline's scoring and matching layer operates entirely on de-identified, aggregated business themes, never on an individual merchant or their personal data.
- **Erasure Scope (open):** on a right-to-erasure request, does deletion cover the raw file only, or every derived row referencing it? Deleting only the file leaves derived data standing; deleting everything downstream retroactively changes scores and strips evidence from a shipped ticket.