

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT

12-Week Hybrid Agentic System Pilot

Monitoring Plan



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · MONITORING PLAN

1. Executive Summary

This plan defines how the Payments pilot's impact is measured against the five success criteria set in the Project Brief.:

- **Reduced time spent on manual synthesis**
- **Greater consistency in product documentation**
- **Improved traceability from customer insight to Jira item**
- **Faster discovery-to-backlog turnaround**
- **Stronger stakeholder confidence in how priorities are set**

Measurement runs on two tracks in parallel:

- **Workflow and business metrics** (PM/Engineering/Support time, ticket quality, backlog visibility)
- **AI system metrics** (gate acceptance, vector matching precision, drafting quality, latency and cost)

Both tracks share the same three-phase pilot structure defined in the Project Brief — Weeks 1–4 (Foundation), Weeks 5–10 (Pilot Execution), Weeks 11–12 (Assessment) — and report through the same Appsmith dashboard and /ask_data agent used elsewhere in the architecture.

2. Monitoring Objectives

- **Baseline first:** Establish a defensible pre-pilot baseline before any AI-assisted step goes live, so gains are measured against real numbers rather than assumption.
- **Track trust, not just speed:** track whether AI-generated outputs are being accepted, edited, or rejected at each of the three human gates, and why.
- **Enable course correction:** surface friction (latency, override rates, error taxonomy) early enough in Weeks 1–6 to tune thresholds and prompts before the Week 11–12 assessment.
- **Support the go/no-go decision:** feed a single Week 11–12 Proceed / Iterate / Stop decision with baseline-vs-actual evidence across both workflow and AI system metrics.

3. Baseline Metrics to Capture (Weeks 1–4)

Baseline data collection runs before or alongside the first live processing, so every later measurement has a pre-AI number to compare against.

Workflow & Business Baselines

Source	Metric	What to Collect
PM	Synthesis time	Hours/week spent reading, summarizing, tagging
PM	Discovery cycle time	Days from feedback received → Jira created
All	Triage meeting duration	Minutes per weekly prioritization meeting, and number of attendees
Engineering	Ticket quality	% of tickets requiring clarification before sprint
Engineering	Jira Ticket Creation Time	baseline duration (minutes) for PMs to write problem statements, user stories, and acceptance criteria
Engineering	Rework rate	How often stories get sent back or rewritten
Support	Ticket volume	Tickets submitted per week/month
Support	Visibility	% of tickets that reached product backlog
All	Confidence	Survey: “Does prioritization reflect real customer needs?” (1–5)
All	Traceability	Can you link a Jira ticket to raw evidence? (yes/no)

AI System Metrics:

- 1) Vector Search & Ingestion Latency** — average latency (ms) for DLP de-identification, embedding generation, and candidate shortlist retrieval
- 2) LLM Token Consumption & Cost** — average tokens and financial cost per processed feedback item and generated draft ticket
- 3) Data DLP Scrubbing Accuracy** — % of raw text successfully anonymized without losing contextual meaning

4. Metrics Framework

Ongoing monitoring is organized into three analytics layers: Feedback Analytics, Process & Workflow Analytics, and AI System Analytics.

4.1 Feedback Analytics

Report	Grouping	Key Metrics & Focus
Emerging Feedback Velocity	Feedback Trends & Velocity	Tracks rolling 7-day and 14-day spikes in raw feedback grouped by product area tag
Feature Area Backlog Depth	Product Area & Regression	Aggregates unresolved opportunities by feature area, prioritized by total feedback volume and severity
Post-Release Regression Tracker	Product Area & Regression	Flags opportunities marked resolved where new feedback continues to arrive post-fix version release date

4.2 Process & Workflow Analytics

Report	Grouping	Key Metrics & Focus
Insight Conversion & Rejection Rate	Lifecycle & Conversion Funnel	Measures proportion of raw insights linked to existing opportunities vs. forming new opportunities vs. rejected
Opportunity Pipeline Lifecycle	Lifecycle & Conversion Funnel	Tracks volume and conversion through stages
Jira Draft Readiness & Sync Lag	Jira Execution & Delivery	Monitors drafts pending approval and measures time elapsed between draft confirmation and Jira creation
Feedback-to-Ticket -Creation Time	Jira Execution & Delivery	Measures days from raw feedback to Jira ticket created — stops at the pilot's own boundary
Feedback-to-Jira Linkage Rate	Traceability	% of created Jira tickets with an unbroken evidence chain back to source
Pending Review Latency & Bottlenecks	Operational Bottlenecks	Calculates average time raw feedbacks, insights, and jira drafts spend in staging queues

4.3 Gate Decision & Review Outcomes

Report	Grouping	Key Metrics & Focus
Gate Acceptance & Override Rates	Gate Decision Outcomes	Tracks proportion of AI suggestions accepted as-is, edited, or rejected across Gate 1, Gate 2, and Gate 3
Story Generation & Point Variance	Drafting Review Outcomes	Measures acceptance of generated user stories/acceptance criteria and delta between AI vs. PM story points
AI Error Taxonomy & Misclassification	Override Root Cause	Breaks down human overrides by feature area, field type (e.g., severity vs. title), and feedback source type
Draft Completeness Rate	Documentation Consistency	% of Gate 3 drafts containing all four required fields (problem statement, user story, acceptance criteria, evidence-linked rationale)

4.4 AI System Analytics

Report	Grouping	Key Metrics & Focus
Similarity Score & Rank 1 Hit Rate	RAG & Matching Precision	Evaluates vector match score distributions and how frequently PMs select the top-ranked candidate
Latency & Token Efficiency	System Cost & Performance	Monitors API response latency, token usage efficiency, and operational costs per item processing

4.5 Compliance Assurance

Report	Grouping	Key Metrics & Focus
DLP Scrubbing Accuracy (Recurring)	Data Privacy	Biweekly sample accuracy vs. the Weeks 1–4 baseline; escalates to Simone Dubois / Legal on any drop below baseline
Works Council Logging Integrity	Employee Monitoring	Periodic audit confirming gate_decisions logs functional role only (e.g., Payment_PM), never individual employee IDs; escalates to Simone Dubois / Legal on any finding

5. 12-Week Monitoring Roadmap

Weeks 1–4 — Foundation

- **Baseline capture:** Capture all workflow/business and AI system baselines listed in Section 3.
- **Instrumentation:** set up DLP logging and complete the DLP latency audit; set the initial vector similarity threshold.
- **Shadow logging:** run the automated ingest in shadow mode alongside existing spreadsheets, without interrupting the team; compare AI-extracted fields (severity, product area, account tier) against what PMs enter manually.
- **Milestone:** $\geq 85\%$ alignment between AI tags and spreadsheet tags before Gate 1 goes live for real decisions.

Weeks 5–10 — Pilot Execution

1. **Ingestion & Gate 1:** PMs review incoming insights only via the Gate 1 interface. Track Gate 1 Acceptance vs. Override Rate. Refine feature-area taxonomy and calibrate similarity thresholds to reduce field-level overrides.

Milestones: ingestion lead time drops from ~20 minutes to < 1 minute per feedback item; RAG hit rate > 80%.

2. **Shortlisting & Gate 2:** automated scoring views replace spreadsheet-based prioritization. Gate 2 becomes a continuous per-item decision inside Appsmith.

Milestone: weekly triage meeting duration reduced by 75%.

3. **Drafting & Gate 3:** automated Jira ticket drafting and sync go live; Track Jira Draft Edit Rate and AI-vs-Reviewer story point variance.

Milestone: Gate 3 story acceptance rate > 85%.


Weeks 11–12 — Assessment

- **Spreadsheet retirement:** total PM time spent managing feedback drops > 80% vs. the Week 1–4 baseline (50% will also be acceptable considering the new SOP and UI).
- **Final measurement:** verify metrics against baseline; publish the full 12-week AI performance and cost/ROI analysis.
- **Decision input:** results measured against baseline feed the end-of-pilot stakeholder survey and the Proceed / Iterate / Stop decision (Brief §8).

6. Reporting Cadence & Ownership

Reporting is delivered continuously through the same dual architecture used elsewhere in the pilot: the Appsmith dashboard provides persistent, real-time views for standard executive monitoring (Product Analytics, Process & Workflow Analytics, AI System Analytics), while the /ask_data MCP agent supports ad-hoc, natural-language queries for any stakeholder who needs a specific cut of the data outside the standard reports.

Picture 1: Click on the diagram to open Feedback Intelligence & AI Analytics portal.



Ownership

Overall accountability for the Monitoring Plan sits with Elena Vance as Head of Product, while day-to-day operation and specialist oversight are distributed as follows:

Area	Owner	Responsibility
Overall Monitoring Plan	Elena Vance (Head of Product)	Accountable for the plan’s continued relevance, the Week 11–12 assessment, and the final Proceed / Iterate / Stop decision.
Day-to-day metric health & dashboard accuracy	Marcus Bauer (Senior PM)	Ensures data is flowing correctly, investigates anomalies, and keeps the operational reports usable for the team.

Area	Owner	Responsibility
AI system metrics (latency, cost, RAG quality, gate acceptance)	External Consultant + Priya Patel (Eng. Lead)	Maintains technical instrumentation and surfaces performance or cost issues.
Compliance & privacy metrics (DLP accuracy, role-only logging integrity)	Simone Dubois (Risk & Compliance)	Owns recurring DLP sample checks and Works Council logging audits; escalates any breach of baseline.
Baseline vs. actual comparison & final report	External Consultant + Elena Vance	Produces the Week 11–12 performance and ROI analysis that feeds the go/no-go decision.

7. Escalation Triggers

Automated BigQuery latency monitoring, with Appsmith/email notification triggers, watches gate turnaround against the 48-hour SLA escalation protocol (Risk Register OP-01), pending Elena Vance's executive agreement on that SLA.

A sustained drop below any Phase-level milestone in Section 5 (e.g., Gate 1 alignment below 85% past Week 4, or RAG hit rate below 80% past Week 6) is treated as a threshold-tuning trigger, prompting taxonomy or prompt adjustments before the next phase begins, rather than waiting for the Week 11–12 assessment.

8. End-of-Pilot Survey & Final Assessment

In Weeks 11–12, measured results are compared against the Weeks 1–4 baseline across both workflow/business and AI system metrics (Brief §8: “measured results against baseline; stakeholder Survey, Proceed / Iterate / Stop decision”).

The stakeholder confidence survey reuses the baseline instrument captured in Section 3 (“Does prioritization reflect real customer needs?”, 1–5) so the end-of-pilot score is directly comparable to the Week 1–4 baseline score.

Outputs feed the single Proceed / Iterate / Stop decision defined in the Project Brief, alongside the full 12-week AI performance and cost/ROI analysis.