

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT

12-Week Hybrid Agentic System Pilot



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026

Executive Summary.....	2
Vocabulary.....	4
Project Brief.....	5
1. Executive Summary.....	6
2. Business Need.....	7
3. Objective & Requirements.....	7
4. Assumptions.....	8
5. Core Workflow Description.....	9
6. Scope of the Pilot.....	14
7. Stakeholder MAP and RACI-matrix.....	15
8. Timeline & Pilot Structure.....	15
9. Success Criteria & Metrics.....	16
10. Key Risks Summary.....	16
11. Governance & Compliance.....	17
Stakeholder Map and RACI-Matrix.....	18
1. 2x2 Stakeholder Influence-Interest Mapping.....	19
2. In-Depth Stakeholder Profiles & Technical System Alignment.....	20
3. Project Scope & RACI Accountability Matrix.....	23
Risk Mitigation Analysis.....	24
1. Executive Summary.....	25
2. Top Risks Heatlist.....	27
3. Executive Decision Requests.....	27
4. Risk Response Plan.....	28
5. Dependencies & Single Points of Failure (SPOF).....	29
6. Compliance and Privacy Risk Notes.....	30
7. Conflicts & Resolutions.....	31
8. Open Questions.....	31
Monitoring Plan.....	32
1. Executive Summary.....	33
2. Monitoring Objectives.....	33
3. Baseline Metrics to Capture (Weeks 1–4).....	34
4. Metrics Framework.....	35
5. 12-Week Monitoring Roadmap.....	37
6. Reporting Cadence & Ownership.....	38
7. Escalation Triggers.....	39
8. End-of-Pilot Survey & Final Assessment.....	39

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · EXECUTIVE SUMMARY

Executive Summary

The Problem

Finexa's product function is strained by high-volume feedback arriving across six disconnected channels. Manual synthesis consumes significant PM time, insights remain fragmented, backlog items lack clear traceability to verified customer needs, and prioritisation still relies heavily on anecdote. Engineering regularly receives incomplete tickets, creating rework. Leadership has requested a controlled pilot to move from fragmented feedback to evidence-based delivery planning without removing PM authority or compromising security and compliance.

Recommended Solution

A **12-week Hybrid Agentic System** pilot in the Payments domain, built on Finexa's existing AI platform and BigQuery as the single source of truth.

Core design principles:

- Upstream Google Cloud DLP de-identifies all data before it reaches the LLM.
- The model performs only cognitive work (atomic feedback splitting, structuring, vector matching of recurring issues, drafting).
- Three mandatory human gates: Gate 1 (confirm reading & grouping), Gate 2 (approve opportunities for drafting), Gate 3 (approve Jira payload & complexity).
- Priority scoring is a transparent, deterministic formula — not model-generated.
- Jira credentials live only in a deterministic wrapper that fires after Gate 3 approval.
- Full audit trail is inherited from BigQuery; Appsmith provides the operational UI; /ask_data enables natural-language querying.

Stakeholder Alignment

- **Elena Vance (Head of Product):** Auditable evidence chain and real-time dashboards for roadmap reviews; clear Proceed/Iterate/Stop decision at Week 12.
- **Marcus Bauer (Senior PM):** Removes bulk of manual synthesis and drafting while retaining full decision authority.
- **Priya Patel (Engineering Lead):** Structural guarantee that nothing reaches the backlog without Gate 3 review.
- **Simone Dubois (Risk & Compliance):** Edge DLP, role-only logging, and full auditability.
- **Liam O'Connor (Support Ops):** Escalations remain visible; regressions surface automatically.

Success Criteria & Monitoring

Measured against five criteria: reduced manual synthesis time, greater documentation consistency, improved insight-to-Jira traceability, faster discovery-to-backlog turnaround, and stronger stakeholder confidence in prioritisation. We have set measurable and ambitious goals:

≥85% Tag Alignment: High accuracy in automated categorization

>80% RAG Hit Rate: Relevant context retrieved successfully consistently

>85% Gate 3 Acceptance Rate: Engineering accepts drafted stories directly

>50% Reduction in Feedback-Management Time: (80% reduction in feedback-management time in the Monitoring Roadmap is treated as an aspirational goal)

Top Risks & Open Decisions

Highest risks: PII leakage, gate rubber-stamping, Works Council exposure, reviewer bandwidth, and infrastructure outage. All have concrete mitigations and owners.

Five decisions currently block live data and require executive action:

1. Simone Dubois – formal approval of Cloud DLP rules and merchant_id GDPR status.
2. Simone / Legal / Works Council – sign-off on role-only logging.
3. Simone / Legal – soft-erasure policy for right-to-erasure requests.
4. Elena Vance – agreement on 48-hour gate SLA and escalation protocol.
5. Simone / Legal – Legitimate Interest sign-off for merchant_id and feedback ingestion.

Compliance

The design treats compliance as a structural constraint rather than an afterthought. All feedback is de-identified by Google Cloud DLP at the edge before any text reaches the LLM.

Gate decisions record only functional roles (e.g. Payment_PM), never individual employee IDs, to satisfy Works Council (BetrVG §87) requirements.

The system is classified as a low-risk internal administrative tool under the EU AI Act. Full auditability is inherited from BigQuery logging.

Open items remaining with Risk & Compliance (DLP rule approval, merchant_id status, soft-erasure policy, and role-only logging sign-off) must be closed before live data ingestion.

Next Step

Close the open decisions in the coming days. Once resolved, the pilot is ready to enter the Foundation phase (system set-up + baseline capture) with a clear runway and measurable checkpoints for the Week 12 go/no-go decision.

Full documentation package (Project Brief, Monitoring Plan, Risk Register, Stakeholder Map & RACI) is available for detailed review.

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · VOCABULARY

Vocabulary

Opportunity — A canonical, validated customer problem or need that has been confirmed at Gate 1 and scored for prioritisation. It can be linked to multiple insights from different sources and is the central object that later becomes a Jira ticket after Gate 3 approval.

Insight — An atomic, structured piece of customer feedback extracted by the LLM from a raw source (interview, support ticket, note, etc.) with a summary, pain points, root-cause indicators, a product-area tag and a provisional theme. One feedback item can produce one or many Insights.

Raw Feedback — The original de-identified text that entered the system after Cloud DLP scrubbing. This text is never seen by the LLM in its original form.

Gate 1 — Mandatory human review of the AI's reading and grouping of Insights. The reviewer confirms the Insight wording and tags, whether it is new or belongs to an existing Opportunity.

Gate 2 — Mandatory human decision on whether a scored Opportunity is worth the effort of drafting a Jira ticket at this moment.

Gate 3 — Final mandatory human review of the drafted problem statement, user story, acceptance criteria and complexity estimate before anything is sent to Jira.

Hybrid Agentic System — Architecture that combines an autonomous LLM reasoning core with deterministic infrastructure (DLP, scoring formula, Jira wrapper) and mandatory human-in-the-loop gates.

Shadow Mode — Period in Weeks 1–4 during which the AI processes feedback in parallel with the existing manual process, but no live decisions are taken from AI outputs. Used to measure alignment and build trust.

RAG (Retrieval-Augmented Generation) — Vector search that surfaces existing Opportunities most similar to a new Insight so the system can suggest linking instead of creating duplicates.

Regression Flag — Automatic signal raised when new customer feedback arrives for an Opportunity that was previously marked resolved in a released Jira Fix Version.


Role-only Logging — Practice of recording only functional roles (e.g. Payment_PM) in gate_decisions — never individual employee IDs — to comply with Works Council (BetrVG §87) requirements.

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · PROJECT BRIEF

12-Week Hybrid Agentic System Pilot

Project Brief



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · PROJECT BRIEF


1. Executive Summary

Finexa's product function is strained by feedback volume arriving across six disconnected channels (CS p. 1). The cost is not only slow synthesis: backlog items are not traceable to verified customer needs, and prioritization runs on anecdote rather than evidence (CS p. 2).

This brief recommends a **12-week pilot in the Payments domain — a Hybrid Agentic System. It combines an autonomous LLM reasoning core with Appsmith UI and workflow automation, deterministic infrastructure, and MCP-driven analytics.**

The Process diagram shows each step and each table in the repository created/updated.

Picture 1: Process Diagram Click on the diagram to open the full view.



Key Architecture Capabilities:

1) Pure Cognitive Focus (Decoupled Infrastructure)

- **Upstream DLP (Non-AI):** De-identification and PII scrubbing run deterministically before text touches the LLM.
- **Core LLM Engine:** Focuses purely on cognitive tasks — atomic 1-to-N feedback deconstruction, ENUM schema extraction, structuring and summarization of feedback, vector search and shortlisting of recurring issues, and structured Jira payload generation in BigQuery staging.
- **Downstream Sync (Non-AI):** Once a human approves Gate 3, a BigQuery wrapper script automatically dispatches approved tickets to Jira.

2) Native Tool & Operational Integration (Appsmith + MCP)

- **Appsmith Operational Hub:** Replaces manual spreadsheets with a unified UI featuring workflow automation and human-in-the-loop Gate 1–3 controls.
- **Vector RAG Matching:** Directly queries BigQuery/Vector databases to surface recurring issues automatically and reads and compares them with the insight.
- **/ask_data MCP Agent:** Connects to BigQuery via Model Context Protocol, allowing to run natural-language operational queries, track override rates, and evaluate product metrics instantly without writing manual SQL or exporting CSVs.

3) Human-in-the-Loop Governance (Gate 1–3 Architecture)

Instead of running as an unmonitored "black box," the system acts as an intelligent copilot that pre-builds structured artifacts for human validation:

- **Gate 1:** Review/approve atomic problem splits, summary, ENUM tags (severity, feature area etc.). Review/approve RAG candidate matches.
- **Gate 2:** Review/approve scoring, approve opportunities for drafting.
- **Gate 3:** Review/approve pre-drafted Jira payloads and estimated complexity.

2. Business Need

Over the last two quarters, Finexa has seen rising customer complaints across payment and dispute workflows alike (CS p. 3) — merchants who don't know whether a payment is pending, failed, or settled, or what a dispute actually requires of them.

The organization already has strong customer access: 25 interviews, 100+ support tickets, ops escalations, strategic account requests, and Jira issues about recurring confusion (CS p. 3). What it lacks is a way to turn that volume into a shared, current, evidence-backed understanding — today, insights live in scattered notes and spreadsheets, prioritization runs on anecdote, and backlog items regularly arrive at engineering incomplete (CS p. 4).

3. Objective & Requirements

Objective: move the Merchant Experience Product Group from fragmented feedback to evidence-based delivery planning (CS p. 4), without adding headcount or introducing unmanaged tooling.

Three conditions, drawn from the stakeholder guardrails below, bound how this is achieved: strict **security and compliance** control, **no unreviewed AI output** reaching the engineering backlog, and **no removal of the PM's decision authority**.

Key Stakeholder Requirements & Guardrails:

- **Elena Vance (Head of Product):** demands fast turnaround and a clear, auditable evidence trail during roadmap reviews (CS p. 6).
- **Marcus Bauer (Senior PM):** wants relief from manual synthesis, but insists the AI stay an assistant, never his replacement (CS p. 7).
- **Priya Patel (Engineering Lead):** will resist adoption unless strict, non-negotiable review checkpoints keep incomplete or hallucinated tickets out of her backlog (CS p. 7).
- **Simone Dubois (Risk & Compliance):** will block the pilot outright without strict security controls, GDPR compliance, and full auditability (CS p. 8).
- **Liam O'Connor (Support Ops):** needs assurance that operational escalations won't be quietly buried under feature requests (CS p. 8).

Key Constraints:

The wider team has already identified several existing constraints any pilot must work within (CS p. 9):

- **Regulated data** — the work involves regulated financial data, handled within strict controls.
- **A short runway** — limited time for implementation.
- **Uneven data quality** — source material varies considerably across interviews, tickets, and notes.
- **Human oversight** — the process depends on human validation at key points.
- **Adoption** — uptake across teams cannot be assumed.

4. Assumptions

Platform & Technical

- Finexa's existing AI platform can host custom reusable commands, and BigQuery is already Finexa's own data platform — not something introduced *for* this pilot.
- Finexa's Jira workspace uses Fix Versions with real release dates. This is what lets a ticket marked "Done" be told apart from one actually shipped — without it, regression detection has no reliable signal to key off.
- Finexa's Jira supports API-based ticket creation via a service account, not manual entry only.

Data sources & Legal

- The six feedback channels are exportable into BigQuery in reasonably clean form, not, say, unstructured audio with no transcription step already in place.
- Merchant account tier (Enterprise/SMB) already exists as a queryable field on Finexa's merchant records.
- A merchant's `merchant_id` is non-personal data for corporate accounts; for sole proprietors it may identify an individual and needs Legitimate Interest sign-off (TBD with Simone).


Operational & Adoption

- Marcus and Priya have real bandwidth to review consistently across the full 12 weeks — the pilot's evidence is only as good as review actually happening.
- Appsmith access is provisioned for both reviewers.

5. Core Workflow Description.

5.1 The Process diagram shows each step and each table in the repository created/updated with four sequential commands — **/process-feedback**, **/log-feedback**, **/prioritize-opportunities**, **/draft-user-story** — plus **/ask-data** running in parallel for read-only queries, all inside Finexa's existing **AI platform**, writing to and reading from **BigQuery** as the SSOT and using **Appsmith** as the workflow orchestrator and UI.

Picture 2: Process Diagram. Click on the diagram to open the full view.



1. **Google DLP (Data Loss Protection)** — scrubbing PII and sensitive financial information downstream prevents payment/merchant data reaching the model.
2. **/process-feedback (Assist)** — splits one transcript or ticket into distinct problems, each with a summary, pain points, root cause, and a provisional theme. Thin material halts and asks for clarification rather than guessing.
3. **/log-feedback (Assist)** — proposes whether this is the same issue as something already confirmed, or something new.
4. **Gate 1 (mandatory at this stage)** — confirms the reading and the grouping. A *verification* task: is this the same theme as something already validated? Because that has a checkable correct answer, a sustained agreement rate is real evidence Gate 1 could eventually move to spot-checks — the only one of the three gates designed to.
5. **/prioritize-opportunities (Augment)** — reads a priority score already computed by a separate, deterministic formula and explains it in plain language. It never calculates the number itself.
6. **Score Opportunity** — five inputs, each pulled from evidence already sitting in the system, no separate manual step. **Customer value** weighs how many Enterprise vs. SMB merchants are affected. **Urgency** combines the single worst severity reported — one "Critical" call from a customer counts, it isn't averaged down by milder reports — with how often the issue recurs. **Strategic alignment** and **operational risk** are read from configuration set ahead of time by Elena's team and engineering, not decided fresh for each item. **Compliance sensitivity** works the same way. Each input carries a tunable weight, and the score is their weighted sum — arithmetic, not a model call. **Implementation complexity** is deliberately not one of the five; it's assigned manually at Gate 3 instead, once a real draft exists to judge, because no reliable signal exists to estimate it at this stage.

7. **Gate 2 (permanent)** — confirms this specific scoring opportunity is worth spending drafting effort on now. A *judgment* task, not a verification one: by this point the item is already a validated, scored opportunity, and deciding what matters right now is shaped by capacity and timing that shift week to week — there's no single correct answer an agreement rate could ever converge toward, so this gate never graduates.
8. **/draft-user-story (Assist)** — converts an approved opportunity into a problem statement, user story, acceptance criteria, and evidence-linked rationale.
9. **Gate 3 (permanent)** — confirms the draft against the original evidence and the suggested implementation complexity.
10. **The deterministic Wrapper Code** — the only component holding Jira credentials — calls Jira's API and posts selected Jira Drafts to Jira.
11. **/ask-data (Augment)** — plain-language, read-only querying of stored insights, running in parallel to the sequence above, any time.

5.2 Capabilities of the Architecture:

- **PII is unreachable by design.** Cloud DLP de-identifies feedback at ingestion, and the AI's service account holds no credential that could retrieve raw source data — it is not instructed to avoid PII, it has no path to it. The residual risk is free-text residue that DLP misses, controlled by pattern validation before ingestion and ongoing sampling.
- **Parsing and structuring with an LLM into a consistent insight schema** and writing to BigQuery, gives both queryability and a foundation for vector search.
- **Vector search** for similar issues + linked opportunities is the highest-leverage part. The same complaint arriving in an interview, a support ticket and a CRM note becomes **one prioritized opportunity carrying three pieces of evidence**, rather than three unconnected backlog entries.
- **Three mandatory human gates sit inside the sequence.** The AI drafts and proposes; a product manager or the engineering lead decides. No command holds Jira credentials, so no AI output can reach engineering unreviewed — a structural guarantee rather than a policy commitment.
- **The audit trail is inherited, not built.** Every gate decision is a stored row and BigQuery logs data access by default — so auditability is a property of the platform Finexa already runs, not custom logging code that must be maintained and trusted to have been switched on.
- **Priority is a formula, not a judgment.** The score is arithmetic over five of the six criteria the case names; implementation complexity is deliberately excluded — it's assigned manually at Gate 3 once a real draft exists, not estimated by the formula. Only the score's explanation is AI-written. Because the weights are parameters, a systematic disagreement at review identifies which weight is wrong and who can change it. An AI-generated score could be disputed, but not tuned.

- **Dedicated UI layer.** People need to see context side-by-side, compare similar opportunities, edit structured fields, and act with one or two clicks.
- **AI Auditability (gate_decisions):** Storing `ai_proposed_value` alongside `pm_final_value` gives you an immediate dataset to measure prompt drift, model accuracy, and human override rates over time.
- **Regression Intelligence (regression_flag + last_insight_at):** Capturing when incoming feedback maps to a resolved or existing canonical opportunity enables automatic detection of post-release regressions or incomplete fixes.

5.3 BigQuery Schema - *Click to open the full view.*

5.4 User Interface

All three gates are executed through a single Appsmith interface rather than across separate tools. Each screen presents the item under review alongside the evidence behind it, with AI-proposed values as editable fields. The three views below show Gate 1, Gate 2, and Gate 3 in sequence.

Human Gate 1:

Picture 4: Insight linked to the source ticket and opportunities. Click on the diagram to open the full view.

The screenshot displays the 'Finexa · Feedback Intelligence' interface. The main content area is divided into two columns. The left column shows an 'Insight #INS-0234' with a 'Support ticket' and 'Enterprise' tag, dated '14 Mar 2025'. It includes a 'Summary' (Merchant cannot see why a payment failed), 'Problem details' (Understand why each payment failed, Recover in the dashboard), 'Jobs & goals' (No failure reason or code, Know what action is in the UI, Support what open internal tools to investigate), 'Pain points' (Failure codes are logged, but not exposed in the merchant-facing payment overview), and a 'Source ticket · ST-49281'. The source ticket details include 'Ticket ID: ST-49281', 'Created: 12 Mar 2025 · 14:22', 'Customer: Acme Payments GmbH (EU · Enterprise)', 'Channel: Email → Support', and an 'Original message (de-identified)' snippet: 'When a payment fails we only see "Payment failed". We need the actual decline reason (insufficient funds, card expired, 3DS failure, etc.) so our team can tell the merchant what to do next. Right now we have to open internal tools for every case.' The ticket has a 'Priority: High' and 'Status: Open'.

The right column shows an 'Opportunity OP-011 · Payment failures lack clear reasons in merchant UI' with a 'Workflow: Payment status' and 'Feature area: payments_status' tag, and a 'for drafting' status. It includes an 'Evidence summary' with a score of '0.82 High priority' and a 'Priority' section with '9 linked insights' and 'Severity: 4 High · 3 Medium · 2 Low' and 'Segments: 4 Enterprise · 5 SMB'. Below this is a 'Recent insights' table:

Insight	Segment	Severity	Excerpt
INS-0210	Enterprise	High	Cannot see why a payme...
INS-0187	Enterprise	High	Need explicit failure reas...

Below the table is a 'Linked Jira tickets (2)' section with two tickets:

- PAY-421 · Expose failure reason in payment details** (In progress, 5 pts). Created 03 Feb 2025 · Assignee: Anna K. Add decline code + human-readable reason to the merchant payment detail view. Include next-best-action suggestion. Tags: payments, merchant-ui.
- PAY-397 · Improve error messaging for failed payments** (Resolved, 3 pts). Created 18 Nov 2024 · Assignee: Tomas R. Replaced generic "Payment failed" with more specific messages for the top 5 decline reasons. Tags: checkout, shipped 2.3.0.

A warning box at the bottom states: 'This opportunity has new insights after a previous fix. Treat as potential regression.'

At the bottom of the left column, 'Matching opportunities' are listed:

- OP-011 · Payment failures lack clear reasons** (0.87). 9 linked insights · High priority · for drafting. View full →
- OP-019 · Generic error messages in checkout** (0.72). 4 linked insights · Medium priority · monitoring. View full →

Human Gate 2:

Picture 5: Opportunity is scored and explained. PM selects opportunities for drafting. Click on the diagram to open the full view.

Finexa · Feedback Intelligence

Gate 2 · Marcus

Prioritized Opportunities

Ranked by score · highest first

#1 OP-011

Regression

Approved

0.87

score

Payment failures lack clear reasons in merchant UI

9 insights · 4 Enterprise · 5 SMB

#2 OP-019

0.74

score

Generic error messages in checkout flow

6 insights · 2 Enterprise · 4 SMB

#3 OP-024

0.61

score

Merchant dashboard lacks payment method health indicators

4 insights · 1 Enterprise · 3 SMB

#4 OP-031

0.53

score

Refund status not visible to merchants in real time

3 insights · 0 Enterprise · 3 SMB

OP-019 · prioritized

Generic error messages in checkout flow

0.74

priority score

Linked insights

6

Severity mix

2H · 3M · 1L

Segments

2 Enterprise · 4 SMB

SCORE EXPLANATION

Recurring feedback from both SMB and enterprise. Lower enterprise weight than OP-011 but still material. Several tickets already linked in previous quarters.

Last scored 5 hours ago

Theme

Checkout error clarity

Workflow

Checkout

Feature area

checkout_flow

Approve for drafting

Skip for now

Human Gate 3:

Picture 6: Jira Draft is checked, complexity is estimated, the draft can be confirmed for Jira. Click on the diagram to open the full view.

Finexa · Feedback Intelligence

Gate 3 · Review Drafts

Jira Drafts

Approved opportunities · ready for review

#1 OP-011

Regression

Confirmed

5 pts

complexity

Expose failure reason in payment details

9 insights · score 0.87

#2 OP-019

3 pts

complexity

Improve error messaging in checkout

6 insights · score 0.74

#3 OP-024

8 pts

complexity

Add payment method health indicators to dashboard

4 insights · score 0.61

LINKED OPPORTUNITY

OP-019

Generic error messages in checkout flow

0.74

score

Checkout error clarity

Checkout

checkout_flow

6 insights

2 Enterprise · 4 SMB

Jira Draft · DRAFT-002

pending_draft

TITLE

Improve error messaging in checkout

PROBLEM STATEMENT

During checkout, customers and merchants see only generic error messages when a payment method is declined. This leads to abandoned checkouts and unnecessary support contacts.

USER STORY

As a customer (or merchant assisting a customer), I want clear, actionable error messages at checkout so I know how to fix the payment problem and complete the purchase.

ACCEPTANCE CRITERIA

Replace generic "Payment failed" with specific messages for the top 5 decline reasons

Messages include a short recovery action (e.g. "Try another card" or "Contact your bank")

Works on both desktop and mobile checkout

COMPLEXITY (STORY POINTS)

1

2

3

5

8

13

AI suggested: 3

Confirm draft & complexity

Edit draft

Send back

12

5.5 /ask_data:

With this volume of clean and structured data continuously flowing through the pipeline, Finexa can dynamically execute three distinct layers of operational analytics: Product Analytics to trace macro feature requirements, Process Analytics to monitor team workflow optimization, and AI Operations Analytics to audit core model efficacy.

To deliver these insights natively, the platform couples an interactive Appsmith frontend interface with a conversational ask_data Analytical Agent powered by the open-standard Model Context Protocol (MCP). This dual architecture allows the organization to seamlessly deploy both persistent, real-time analytics dashboards for standard executive monitoring and ad-hoc, natural-language reporting for immediate, flexible data exploration.

Picture 7: Click [/ASK DATA](#) to open the full report prototype.



6. Scope of the Pilot

In-Scope (Phase V1 — 12 Weeks)

- **Project Documentation & Governance** — produce Project Brief, Stakeholder Map + RACI, Risk Register, Monitoring Plan. Create/update an internal AI application register tracking.
- **Technical Documentation:** define requirements, produce architecture/schema/command documentation
- **Baseline & measurement:** baseline data report + PM/Eng/Support/stakeholder measures + targets
- **Five AI commands:** confirm requirements/guardrails, configure/test, assess quality + adoption
- **BigQuery schema:** confirm data/governance requirements, create tables/relationships/status fields
- **Config tables:** agree criteria/weights/policy rules, create scoring + taxonomy configs
- **Appsmith UI:** define user journeys/approval roles, build command + review screens
- **Human gates:** agree responsibilities, implement Gate 1/2/3 workflows with decision logs
- **Jira wrapper + sync:** confirm access boundaries, build credential-isolated wrapper, status sync
- **Source data ingestion:** confirm baseline dataset, load + de-identify interviews and tickets
- **Monitoring plan (pilot-level):** define metrics/collection, instrument logs + measurement points. Final project success assessment against the baseline.

Out of scope (V1) — Deferred to enterprise rollout (V2+)

- Enterprise AI governance expansion (standing committees, org-wide AI inventory/register)
- Fine-tuning / continual learning loops (training pipelines, RLHF, automated retraining)
- Enterprise monitoring platform + fully formal 24/7 incident operations
- Autonomous orchestration/triggers (e.g., n8n)
- Full channel coverage / org-wide rollout beyond the pilot domain and dataset
- Executive dashboards / Slack bots beyond what's required for pilot evaluation
- Deep UI productization (full UX research, accessibility/localization hardening)
- Organization-wide change management (company-wide training + adoption KPIs; HR/performance implications beyond pilot-safe logging)

The deferred capabilities outlined in the 'Out of Scope' section serve as the foundation for the Phase V2 Enterprise Rollout, which will be considered upon the successful validation of the Phase V1 pilot metrics against the historical baseline. Once this success threshold is reached and executive approval is secured, the current operational boundaries can expand, shifting all V1 exclusions directly into active, in-scope delivery requirements.

7. Stakeholder MAP and RACI-matrix.

Full details in the Stakeholder MAP and RACI-matrix.

Picture 4: Stakeholder Map and RACI Matrix. Click on the diagram to open the full view.

01 / STAKEHOLDER DIRECTORY		02 / RACI MATRIX		
STAKEHOLDER	POWER	INTEREST	STRATEGY	KEY PRIORITIES & GOVERNANCE FOCUS
Elena Vance HEAD OF PRODUCT	HIGH	HIGH	MANAGE CLOSELY	ROI, overall business value, roadmap alignment, scope approval, and governance risk.
Marcus Bauer SENIOR PM	MEDIUM	HIGH	MANAGE CLOSELY	Product specifications, command accuracy, prompt scope, and target baseline metrics.
Priya Patel ENG LEAD	HIGH	HIGH	MANAGE CLOSELY	System architecture, BigQuery schema, technical debt, UI performance, and integration reliability.
Simone Dubois RISK & COMPLIANCE	HIGH	MEDIUM	KEEP SATISFIED	PII de-identification, Cloud DLP compliance, audit decision logs, and regulatory boundaries.
Liam O'Connor SUPPORT OPS	MEDIUM	HIGH	KEEP INFORMED	Operational usability, source data ingestion accuracy, and end-user support workflow efficiency.
External Consultant IMPLEMENTATION PARTNER	LOW	HIGH	KEEP INFORMED	Clear technical scope, rapid issue unblocking, pipeline execution, and sprint velocity.

8. Timeline & Pilot Structure

Picture 5: Project Implementation Roadmap. Click on the diagram to open the full view.



Weeks 1–4 (Foundation): DLP template sign-off, schema build, baseline capture, stakeholder survey, stakeholder RACI agreement, Appsmith configuration, processing of the existing payments data.

Weeks 5–10 (Pilot Execution): live processing of daily Payments feedback.

Weeks 11–12 (Assessment): measured results against baseline; stakeholder confidence Survey, Proceed / Iterate / Stop decision.

9. Success Criteria & Metrics

Full details in the Monitoring Plan

- A measurable reduction in time spent on manual synthesis.
- Greater consistency in product documentation.
- Improved traceability from a customer insight through to the Jira item it informs.
- A faster turnaround from discovery to backlog readiness.
- Stronger stakeholder confidence in how priorities are set (CS p. 9).

10. Key Risks Summary

These are the risks stakeholders raised directly (CS p. 6–8).

Risk	What could go wrong	Control / safeguard
Backlog flooding / low-quality tickets	Engineering receives many inaccurate/incomplete Jira items	Human approval before “ready for backlog”
Over Reliance on AI	Teams accept drafts without critical thinking	Mandatory evidence citations, human sign-off
Privacy/confidentiality leakage	Direct identifiers or unnecessary financial details enter system	Sanitize before ingestion, store only customer reference
Bias in synthesis/prioritization	Important ops/compliance issues are deprioritized	Explainable scoring that can be tuned + human approval
Adoption stalls	Teams ignore the process, return to spreadsheets	Friendly user interface, clear SOP, fast visible wins

The full analysis and resulting top-5 risks ranking are in the Project's Risks Hub.

Picture 5: Finexa Pilot Executive Risk & Activity Hub. Click to review.

Risk Cockpit Risk Resolution Activity Hub Executive Analytics Taxonomy & Scoring Framework		
🔥 Top Risks Headlist		
ID	Risk Title	Severity
RE-01	Data Privacy & PII Leakage	20
TE-02	Human Gate Rubber-Stamping / Cognitive Fatigue	16
RE-02	Works Council Exposure (BetrVG §87)	15
OP-01	Operational Bandwidth & Gate Latency	15
TE-06	Total Infrastr./ API Outage	15

11. Governance & Compliance

- **Regulated Financial Data:** protected via Cloud DLP de-identification, EU-only data residency, and a defined retention policy on raw files (CS p. 9).
- **EU AI Act Alignment:** classified as an internal drafting/administrative tool, falling outside the high-risk tier — it makes no autonomous decisions affecting an individual's credit, employment, or rights.
- **GDPR Compliance:** relies on Legitimate Interest for feedback ingestion and merchant_id resolution, pending sign-off. merchant_id stays plain, not tokenized, and never travels past the point where the account tier is resolved at Gate 1; every table a PM or the AI can see downstream carries the tier only, never the merchant's identity. Individual names, emails, and phone numbers within free text are what Cloud DLP redacts, before any structured table exists.
- **Works Council (BetrVG §87):** decision logs capture reviewer **role** (e.g., Payment_PM), not individual employee IDs, to avoid performance-monitoring concerns.
- **Data Protection Impact Assessment:** current assessment is none required (pending confirmation). DPIA-relevant triggers concern processing *about natural persons* — scoring, profiling, evaluation. This pipeline's scoring and matching layer operates entirely on de-identified, aggregated business themes, never on an individual merchant or their personal data.
- **Erasure Scope (open):** on a right-to-erasure request, does deletion cover the raw file only, or every derived row referencing it? Deleting only the file leaves derived data standing; deleting everything downstream retroactively changes scores and strips evidence from a shipped ticket.

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · STAKEHOLDERS

12-Week Hybrid Agentic System Pilot

Stakeholder Map and RACI-Matrix



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026


From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · STAKEHOLDERS

1. 2x2 Stakeholder Influence-Interest Mapping

The organization is mapped across classic power and interest quadrants to ensure that system-level technical parameters directly satisfy individual constraints and mitigate delivery risks

Picture 1: Stakeholder Map. Click on the diagram to get to the Stakeholder Directory.



The project's key stakeholders show exceptionally high interest in this project and strong motivation to adopt this AI innovation to eliminate repetitive documentation work and drive evidence-based planning.

Team leadership explicitly recognizes the urgent need for a centralized, automated repository to resolve fragmented workflows and communication lines. This strong cross-functional engagement ensures active, collaborative participation throughout the human-in-the-loop pilot phase to maximize long-term business value.

As the executive champion who launched this initiative, Elena Vance provides critical sponsorship and high enthusiasm for the pilot's potential. Her active advocacy ensures strong strategic backing, as she explicitly relies on this system to establish the auditable data trails and product decision traceability required for high-level roadmap reviews.

2. In-Depth Stakeholder Profiles & Technical System Alignment

Before we begin reviewing individual stakeholder criteria and crossing out specific risk guardrails it is critical to address the foundational philosophy behind the architecture we recommend. Every core operational bottleneck, personal friction point, and systemic hesitation was meticulously analyzed before the final process design was confirmed.

The system is engineered to technically enforce best practices directly within the underlying data layer. By utilising native GSC functions, using deterministic constraints — such as edge tokenization, strict database relationships, taxonomy, and isolated dispatch wrapper — we have built a system that seamlessly guides users toward compliant, seamless execution.

Ultimately, this architecture ensures that the team is structurally aligned to achieve maximum velocity and full auditability without adding a single minute of manual administrative overhead.

Manage Closely (High Power / High Interest)

Elena Vance (Head of Product): Executive Champion

Systemic Frustration: Highly frustrated that product prioritization discussions rely heavily on anecdotal evidence and fragmented spreadsheets. This manual fragmentation makes it impossible to establish a clear, auditable link to verified customer needs during high-stakes roadmap reviews. Elena's main pressure point is time. She operates on a highly limited pilot runway and worries that team adoption friction will delay business value realization.

AI Motivation & Stance: Served as the executive champion who launched the AI pilot initiative. She is highly enthusiastic about reducing low-value manual work, increasing product decision traceability, and driving measurable business value across complex domains.

Hesitation & Risk Guardrails: Main pressure point is execution speed. Elena operates on a highly limited timeline for pilot implementation and worries that adoption challenges across internal teams could slow down the realization of these core benefits.

Technical Architecture Alignment: The insight repository database of raw feedbacks, processed insights, opportunities and jira drafts as a SOT. Elena can generate any Product, Process, or AI operations report instantly through the Appsmith dashboard or via simple, natural-language queries using the /ask_data agent. Absolute traceability to customer data and zero manual reporting effort. Team onboarding isn't assumed — it is monitored and controlled (See. Monitoring Plan).

Priya Patel (Engineering Lead — Merchant Experience)

Systemic Frustration: Distressed that her development team regularly receives incomplete Jira tickets and loosely defined requirements. This directly triggers constant, wasteful clarification loops, engineering rework, and severely delayed delivery cycles.

AI Motivation & Stance: Highly skeptical of the AI rollout. Priya fears that product managers will use the automated tools to rapidly flood her engineering backlog with inaccurate, unverified, or incomplete AI-generated user stories. She is structurally resistant to a general overreliance on these outputs without deep human critical thinking.

Strict Non-Negotiable Demand: Priya will actively resist the tool's adoption unless strict, non-negotiable human review checkpoints are permanently established before *any* candidate Jira ticket is pushed to engineering pipelines.

Technical Architecture Alignment: Handled by the deterministic Jira Wrapper API isolation boundary. The automated system lacks deployment credentials and cannot bypass human evaluation. Gate 3 forces a mandatory draft layout validation and manual story-point configuration by engineering, ensuring zero backlog flooding.

Marcus Bauer (Senior Product Manager)

Systemic Frustration: Exhausted by spending hours manually reading, summarizing, and tagging themes across dozens of merchant interview transcripts and support tickets instead of doing strategic product work.

AI Motivation & Stance: Positioned as an eager early adopter. Marcus is keen to use reusable AI commands to automate tedious documentation and speed up the turnaround from discovery to backlog readiness.

Hesitation & Risk Guardrails: Quietly worries about overreliance on generated outputs. Marcus requires firm reassurance that the AI will act strictly as an assistant rather than replacing his ultimate product decision authority.

Technical Architecture Alignment: Handled via the Appsmith Operational UI Hub. The interface streamlines documentation tasks while maintaining manual override dropdown controls across Gates 1 and 2. To prevent automation complacency, the text layout splits descriptions into a rigid Known (Fact + Direct Text Citation) vs Assumptions (Model Interpretations) system, ensuring Marcus remains the absolute decision authority.

Keep Satisfied (High Power / Medium Interest)

Simone Dubois (Risk and Compliance Manager)

- **Systemic Frustration:** Strained by current inconsistent handoffs between product discovery and delivery, which makes verifying that strict compliance and risk guidelines are met an uphill battle.
- **AI Motivation & Stance:** Represents the strongest layer of organizational resistance due to the strict guardrails surrounding regulated financial data handling. Simone is deeply concerned about financial data privacy, confidentiality leakage, and the potential for structural bias in AI-driven prioritization algorithms.
- **Strict Non-Negotiable Demand:** Simone will block the pilot entirely if she detects a lack of data security controls. She demands that all AI outputs remain fully auditable, secure, and reliable, with a clear human-in-the-loop validation process.

Technical Architecture Alignment: Addressed through the **Google Cloud DLP Pipeline**. Financial details and customer PII are tokenized at the ingestion point. The system relies on Legitimate Interest for feedback analysis (pending sign-off), and enforces strict **BetrVG §87** compliance by recording reviewer interactions in `gate_decisions` solely by functional corporate role. Handoff between discovery and delivery is the opportunity status change, done by PM - Marcus, when he confirms the opportunity is ready for drafting.

Keep Informed (Medium-Low Power / High Interest)

Liam O'Connor (Customer Support Operations Manager)

Systemic Frustration: Deeply frustrated that despite passing along 100+ support tickets and operational escalations regarding merchants' confusion over payment and balance statuses, these insights remain completely trapped in separate tools and disconnected from the actual product backlog.

AI Motivation & Stance: Highly optimistic about the feedback-logging step. Liam believes it could create a centralized, searchable customer insights repository that finally gives his support tickets a permanent, visible influence in the core product lifecycle.

Hesitation & Risk Guardrails: More resistant when it comes to using AI for the final prioritization of opportunities. He fears the system might suffer from synthesis bias, inadvertently deprioritizing crucial operational escalations in favor of purely feature-driven requests.

Technical Architecture Alignment: Addressed by the system `regression_flag` and `last_insight_at` markers. When incoming customer feedback maps dynamically to a pre-existing opportunity previously marked as resolved in Jira, the system triggers a regression warning. This forces the issue back into the active review queue, ensuring operational escalations maintain visibility and cannot be quietly buried by features.

3. Project Scope & RACI Accountability Matrix

To establish operational clarity before project handover, the pilot provides an interactive **RACI Governance Hub** directly within the deployment dashboard. This interface functions as a live accountability registry where internal team members can filter by their individual name profiles.

Picture 2: Project RACI Matrix. Click on the diagram and choose second tab to get more information.

The screenshot displays the '02 / RACI MATRIX' tab. At the top, six individual profiles are shown with their roles and RACI counts:

Profile	Role	Accountable	Responsible	Consulted	Informed
Elena Vance	HEAD OF PRODUCT	3	0	1	8
Marcus Bauer	SENIOR PM	3	1	4	4
Priya Patel	ENG LEAD	3	1	4	4
Simone Dubois	RISK & COMPLIANCE	2	0	4	6
Liam O'Connor	SUPPORT OPS	1	0	2	9
External	CONSULTANT	0	10	2	0

Below the profiles is a filter bar and a dropdown menu set to 'Elena Vance (Head of Product)'. The main table lists deliverables and their RACI assignments across the team members:

DELIVERABLE & FIELD SPECS	ELENA	MARCUS	PRIYA	SIMONE	LIAM	EXTERNAL
Project & Governance Docs Project Brief, Stakeholders, Risk Register, & AI App Register → app_register	A	C	I	C	I	R
Baseline & Measurement Metrics collection (PM/Eng/Support) & baseline targets report → baseline_metrics	A	R	C	I	C	C
Monitoring Plan & Assessment Metrics logging, log instrumentation, & final success assessment → monitoring_logs	A	I	I	I	C	R

Qualifications-Based Accountability Distribution and Process Ownership

Accountability (A) across the 12 core pipeline components has been evenly balanced and assigned to the leadership group. This mapping guarantees that process ownership matches individual technical qualifications and project interests:

- **Elena Vance (Head of Product):** Governs high-level strategic alignment, business target measurements, and value tracking views
- **Marcus Bauer (Senior PM):** Directs daily interface touchpoints, user journey reviews, and administrative automation features
- **Priya Patel (Engineering Lead):** Safeguards backend engineering pipelines, database schemas, and integration access controls
- **Simone Dubois (Risk & Compliance):** Controls the deterministic data sanitization pipeline and system action logging configurations
- **Liam O'Connor (Support Ops Manager):** Manages cross-channel incoming feedback conduits to protect customer escalation pathways

Assigning ownership at the outset keeps these owners in the loop throughout the build, so the transition to steady-state operation carries no knowledge gap — they have been accountable for their component from the beginning, not handed it at the end.

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT

12-Week Hybrid Agentic System Pilot

Risk Mitigation Analysis



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026


From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · RISK MITIGATION ANALYSIS

1. Executive Summary

To ensure an easy overview of the Project's risks and required actions and decisions we have published the entire Project risk register into a live, interactive **Risk Management Hub**.

Picture 1: Finexa Executive Risk Hub, Taxonomy & Scoring Framework. Click on the diagram to get more information.



Risks Requiring Decisions Now

- **GDPR Right-to-Erasure Policy:** Formal executive decision on whether a right-to-erasure request triggers a soft erasure (anonymizing derived insights) or a hard cascade wipe (purging historical scores and ticket lineage)
- **Works Council (BetrVG §87) Sign-Off:** Explicit agreement on reviewer anonymity protocols (logging functional roles like Payment_PM rather than employee IDs) before live execution (Pilot Project Brief, p. 11).
- **Legitimate Interest & merchant_id Classification:** Formal LIA sign-off and approval of the entity-type classification rule (sole proprietor vs. corporate) governing merchant_id and feedback-ingestion processing.

Top 5 Risks

1. **RE-01 (Severity 20):** Incomplete Cloud DLP redaction leaks merchant PII or financial data into LLM prompts, leading to regulatory sanctions under GDPR and immediate pilot blockage by Risk & Compliance.
2. **TE-02 / OP-05 (Severity 16):** Reviewer cognitive fatigue causes human gatekeepers to blindly rubber-stamp AI drafts, leading to low-quality/hallucinated user stories entering the engineering backlog.
3. **RE-02 (Severity 15):** Logged review actions enable indirect performance tracking of PMs, leading to Works Council (BetrVG §87) intervention and project suspension.
4. **OP-01 (Severity 15):** Daily PM operational duties delay human gate approvals beyond 48 hours, leading to workflow bottlenecks and schedule slippage across the 12-week pilot runway.
5. **TE-06 (Severity 15):** Upstream LLM provider outages or API downtime halt the Appsmith UI, leading to total operational paralysis of the discovery pipeline.

Key Assumptions Affecting Risk Posture

- **Merchant Account Tier availability:** merchant_account_tier exists as a queryable field in Finexa merchant databases.
- **Merchant ID classification and Legitimate Interest basis are pending Simone Dubois/Legal sign-off.**
- **Core reviewers (Marcus Bauer, Priya Patel) have dedicated bandwidth (~15–30 min daily) throughout all 12 weeks.**
- **BigQuery and Jira service accounts, including REST API ticket creation,** can be safely integrated with credential isolation.

Biggest Exposure Themes

- **Compliance & Data Privacy:** Exposure around GDPR data handling, EU AI Act categorization, and Works Council employee monitoring constraints (BetrVG §87)
- **Human Gate Operations & Adoption:** Dependencies on Marcus Bauer and Priya Patel maintaining consistent review bandwidth without succumbing to rubber-stamping or bypassing Appsmith for spreadsheets (CS, p. 6–7).
- **Engineering Backlog Quality:** Operational friction caused if unverified, poorly structured, or hallucinated requirements bypass Gate 3 and reach engineering sprints.
- **System Integration & Failure Recovery:** Potential failure points across BigQuery schema changes, Jira API wrapper timeouts or prompt injection attacks.

2. Top Risks Heatlist

ID	Risk Title	Severity	Next Important Action	Decision Needed
RE-01	Data Privacy & PII Leakage	20	Validate Cloud DLP regex patterns against raw sample dataset in Week 1 sandbox	Simone Dubois formal approval of Cloud DLP de-identification rules before live data ingestion
TE-02	Human Gate Rubber-Stamping / Cognitive Fatigue	16	Configure Appsmith UI to enforce mandatory rationale text entry on Gate 3 approvals and lock 20% random sample for manual complexity check	Marcus Bauer & Priya Patel agreement on mandatory gate fields and audit thresholds
RE-02	Works Council Exposure (BetrVG §87)	15	Strip individual user identifiers from gate_decisions schema and verify role-only logging (Payment_PM)	Formal sign-off from Works Council representative / Legal on anonymized logging design
OP-01	Operational Bandwidth & Gate Latency	15	Deploy automated BigQuery latency monitoring view with Appsmith/email notification triggers	Executive agreement from Elena Vance on 48-hour SLA escalation protocol
TE-06	Total Infrastr./ API Outage	15	Build and test the BYPASS_AI = TRUE manual CRUD toggle in Appsmith UI	Engineering Lead sign-off on fallback interface readiness

3. Executive Decision Requests

Decision Needed	Decision Owner	Target	Impact if Delayed
Approval of Cloud DLP de-identification rules & merchant_id handling	Simone Dubois (Risk & Compliance)	Week 1	Live data ingestion cannot begin; pilot halted.

Decision Needed	Decision Owner	Target	Impact if Delayed
Adoption of Soft Erasure Policy for GDPR right-to-erasure requests	Simone Dubois / Legal	Week 2	Risks data loss or non-compliance during pilot execution.
Sign-off on Works Council anonymity framework (role vs user_id)	Simone Dubois / HR / Works Council	Week 2	Risk of labor union objection and system shutdown.
Executive mandate for 48-hour gate latency SLA and escalation protocol	Elena Vance (Head of Product)	Week 2	Unresolved review bottlenecks cause pilot schedule overrun.

4. Risk Response Plan

Next 10 Actions in the Next 2 Weeks

Action / Task	Description	Owner	Mitigated Risk
Execute Cloud DLP Validation Test	Run DLP scrubbing on 25 sample interview transcripts to verify 0% PII leak rate.	Simone Dubois	RE-01
Configure Role-Only Logging Schema	Update gate_decisions BigQuery table DDL to log functional role metadata only.	External Consultant	RE-02
Establish Baseline Metric Capture	Finalize pure manual baseline timing benchmarks in BigQuery	External Consultant	OP-04
Provision Appsmith Access & Roles	Grant Appsmith UI access and permission roles to Marcus Bauer and Priya Patel.	Priya Patel	OP-02 & OP-01
Configure Latency Alert Telemetry	Build automated query alerting Elena Vance if avg_gate_latency_hours > 48.	External Consultant	OP-01

Action / Task	Description	Owner	Mitigated Risk
Implement Deterministic Priority Formula View	Deploy opportunity_priority_metrics SQL view with weighted formula.	External Consultant	TE-01
Build Jira Wrapper Two-Phase Commit	Code PENDING_SYNC status step in Jira dispatch wrapper.	Priya Patel	TE-07
Establish Read-Only MCP Connection	Configure /ask_data Model Context Protocol agent with read-only IAM rights.	Priya Patel	TE-08
Set Up Gate 3 Forced Sample Check	Code Appsmith logic locking every 5th draft for manual complexity entry.	External Consultant	TE-02
Publish Command Execution SOP	Deliver standard operating guidelines for /process-feedback and /draft-user-story.	Marcus Bauer	TE-03
Confirm legal_entity_type classification rule	Verify legal_entity_type exists on merchant KYC records; define sole-proprietor vs. corporate classification rule for merchant_id.	Simone Dubois	RE-01 / RE-02

5. Dependencies & Single Points of Failure (SPOF)

SPOF	Why It Matters	Risk(s)	Owner	Mitigation
Google Cloud DLP Engine	Essential edge control preventing raw PII from reaching third-party LLM APIs (Pilot Brief, p. 5).	RE-01	Simone Dubois	Implement strict pre-ingestion validation checks and automated regex tests.
BigQuery Central Staging Repository	Serves as the Single Source of Truth (SSOT) for feedback, vector embeddings, and gate decisions (Pilot Brief, p. 5).	TE-06, TE-10	Priya Patel	Enforce schema abstraction views and automated daily BigQuery backups.

SPOF	Why It Matters	Risk(s)	Owner	Mitigation
Jira Dispatch Wrapper Script	Sole component holding Jira API credentials; bridges BigQuery to engineering backlog (Pilot Brief, p. 5).	TE-06, TE-07	Priya Patel	Two-Phase Commit architecture (PENDING_SYNC) and isolated service account key management.
Marcus Bauer (Senior PM Bandwidth)	Primary human gatekeeper for Gates 1 and 2 in the Payments domain (CS, p. 6).	OP-01, OP-02	Elena Vance	Appsmith UI workflow compression (88% time savings) and backup reviewer designation.

6. Compliance and Privacy Risk Notes

Applicable Policies & Regulations

- **GDPR (EU General Data Protection Regulation):** Applies to merchant feedback handling, data residency (EU-only), and right-to-erasure enforcement (Pilot Brief, p. 11).
- **EU AI Act:** Classified as an internal administrative/drafting tool (low risk), falling outside high-risk regulatory constraints as it makes no autonomous binding decisions (Pilot Brief, p. 11).
- **German Works Constitution Act (BetrVG §87):** Regulates employee monitoring and co-determination rights regarding performance logging (Pilot Brief, p. 11).

Project-Specific Compliance Risk Statements

- **RE-01 (PII Ingestion):** If unstructured customer transcripts contain unscrubbed payment card numbers or personal emails, then passing them to the LLM core violates GDPR data protection mandates, leading to regulatory fines and project shutdown.
- **RE-02 (Employee Monitoring):** If gate_decisions tables log individual PM timestamps and review speeds, then the system constitutes performance monitoring under BetrVG §87, leading to Works Council injunctions.

Open Compliance Questions

- Does the formal right-to-erasure policy permit Soft Erasure (anonymizing derived insights) without requiring cascading deletion across historical Jira tickets?
- Is merchant_id explicitly re-confirmed as non-personal corporate data by Legal when company names are retained for tier resolution at Gate 1?

7. Conflicts & Resolutions

Conflict Discovered	Originating Sources	Impact on Risk Posture	Owner	Proposed Resolution
Speed vs. Governance Checkpoints	Elena Vance demands fast turnaround (CS, p. 5); Priya Patel demands strict, non-negotiable review checkpoints (CS, p. 6).	Friction over review speed; potential for gate bypass or bottlenecking.	Elena Vance	Maintain strict Gate 1–3 checkpoints, but optimize UI layout to compress review time to < 4 minutes/item.
Prioritization Automation vs. Bias Concerns	Executive vision promotes automated priority scoring (CS, p. 2); Liam O'Connor fears operational escalations will be buried (CS, p. 7).	Stakeholder resistance and lack of trust in roadmap outputs.	Elena Vance / Marcus Bauer	Use deterministic formula where operational risk carries explicit tunable weight, combined with automated regression_flag override.
Erase Scope (Hard Cascade vs. Soft Anonymization)	Compliance notes flag unresolved choice between deleting raw files vs. cascade deleting derived rows (Pilot Brief, p. 11).	Data loss vs. compliance	Simone Dubois	Adopt Soft Erasure Policy: purge direct identifiers while preserving anonymized derived insights for backlog integrity.

8. Open Questions

Resourcing & Ownership: Who acts as the official backup gate reviewer when Marcus Bauer or Priya Patel are on leave?

Vendors / Procurement: What are the rate limits and SLA guarantees on Finexa's current AI platform tenant?

Technical & Data Constraints: Are raw interview audio files being transcribed prior to ingestion, or is transcription within pilot scope?

Compliance & Security: Has a formal DPIA exemption letter been signed off?

Stakeholder Alignment & Adoption: What specific target score on the end-of-pilot stakeholder survey defines "stronger stakeholder confidence"?

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT

12-Week Hybrid Agentic System Pilot

Monitoring Plan



Prepared by
Margarita Sergeeva
AI Project Management | THRIVE & Kiron Education
margaret.sergeeva@gmail.com
www.linkedin.com/in/margaritasergeeva/

August 2026

From Customer Voice to Delivery Execution

FINEXA · MERCHANT EXPERIENCE PRODUCT GROUP · PAYMENTS PILOT · MONITORING PLAN

1. Executive Summary

This plan defines how the Payments pilot's impact is measured against the five success criteria set in the Project Brief.:

- **Reduced time spent on manual synthesis**
- **Greater consistency in product documentation**
- **Improved traceability from customer insight to Jira item**
- **Faster discovery-to-backlog turnaround**
- **Stronger stakeholder confidence in how priorities are set**

Measurement runs on two tracks in parallel:

- **Workflow and business metrics** (PM/Engineering/Support time, ticket quality, backlog visibility)
- **AI system metrics** (gate acceptance, vector matching precision, drafting quality, latency and cost)

Both tracks share the same three-phase pilot structure defined in the Project Brief — Weeks 1–4 (Foundation), Weeks 5–10 (Pilot Execution), Weeks 11–12 (Assessment) — and report through the same Appsmith dashboard and /ask_data agent used elsewhere in the architecture.

2. Monitoring Objectives

- **Baseline first:** Establish a defensible pre-pilot baseline before any AI-assisted step goes live, so gains are measured against real numbers rather than assumption.
- **Track trust, not just speed:** track whether AI-generated outputs are being accepted, edited, or rejected at each of the three human gates, and why.
- **Enable course correction:** surface friction (latency, override rates, error taxonomy) early enough in Weeks 1–6 to tune thresholds and prompts before the Week 11–12 assessment.
- **Support the go/no-go decision:** feed a single Week 11–12 Proceed / Iterate / Stop decision with baseline-vs-actual evidence across both workflow and AI system metrics.

3. Baseline Metrics to Capture (Weeks 1–4)

Baseline data collection runs before or alongside the first live processing, so every later measurement has a pre-AI number to compare against.

Workflow & Business Baselines

Source	Metric	What to Collect
PM	Synthesis time	Hours/week spent reading, summarizing, tagging
PM	Discovery cycle time	Days from feedback received → Jira created
All	Triage meeting duration	Minutes per weekly prioritization meeting, and number of attendees
Engineering	Ticket quality	% of tickets requiring clarification before sprint
Engineering	Jira Ticket Creation Time	baseline duration (minutes) for PMs to write problem statements, user stories, and acceptance criteria
Engineering	Rework rate	How often stories get sent back or rewritten
Support	Ticket volume	Tickets submitted per week/month
Support	Visibility	% of tickets that reached product backlog
All	Confidence	Survey: “Does prioritization reflect real customer needs?” (1–5)
All	Traceability	Can you link a Jira ticket to raw evidence? (yes/no)

AI System Metrics:

- 1) Vector Search & Ingestion Latency** — average latency (ms) for DLP de-identification, embedding generation, and candidate shortlist retrieval
- 2) LLM Token Consumption & Cost** — average tokens and financial cost per processed feedback item and generated draft ticket
- 3) Data DLP Scrubbing Accuracy** — % of raw text successfully anonymized without losing contextual meaning

4. Metrics Framework

Ongoing monitoring is organized into three analytics layers: Feedback Analytics, Process & Workflow Analytics, and AI System Analytics.

4.1 Feedback Analytics

Report	Grouping	Key Metrics & Focus
Emerging Feedback Velocity	Feedback Trends & Velocity	Tracks rolling 7-day and 14-day spikes in raw feedback grouped by product area tag
Feature Area Backlog Depth	Product Area & Regression	Aggregates unresolved opportunities by feature area, prioritized by total feedback volume and severity
Post-Release Regression Tracker	Product Area & Regression	Flags opportunities marked resolved where new feedback continues to arrive post-fix version release date

4.2 Process & Workflow Analytics

Report	Grouping	Key Metrics & Focus
Insight Conversion & Rejection Rate	Lifecycle & Conversion Funnel	Measures proportion of raw insights linked to existing opportunities vs. forming new opportunities vs. rejected
Opportunity Pipeline Lifecycle	Lifecycle & Conversion Funnel	Tracks volume and conversion through stages
Jira Draft Readiness & Sync Lag	Jira Execution & Delivery	Monitors drafts pending approval and measures time elapsed between draft confirmation and Jira creation
Feedback-to-Ticket -Creation Time	Jira Execution & Delivery	Measures days from raw feedback to Jira ticket created — stops at the pilot's own boundary
Feedback-to-Jira Linkage Rate	Traceability	% of created Jira tickets with an unbroken evidence chain back to source
Pending Review Latency & Bottlenecks	Operational Bottlenecks	Calculates average time raw feedbacks, insights, and jira drafts spend in staging queues

4.3 Gate Decision & Review Outcomes

Report	Grouping	Key Metrics & Focus
Gate Acceptance & Override Rates	Gate Decision Outcomes	Tracks proportion of AI suggestions accepted as-is, edited, or rejected across Gate 1, Gate 2, and Gate 3
Story Generation & Point Variance	Drafting Review Outcomes	Measures acceptance of generated user stories/acceptance criteria and delta between AI vs. PM story points
AI Error Taxonomy & Misclassification	Override Root Cause	Breaks down human overrides by feature area, field type (e.g., severity vs. title), and feedback source type
Draft Completeness Rate	Documentation Consistency	% of Gate 3 drafts containing all four required fields (problem statement, user story, acceptance criteria, evidence-linked rationale)

4.4 AI System Analytics

Report	Grouping	Key Metrics & Focus
Similarity Score & Rank 1 Hit Rate	RAG & Matching Precision	Evaluates vector match score distributions and how frequently PMs select the top-ranked candidate
Latency & Token Efficiency	System Cost & Performance	Monitors API response latency, token usage efficiency, and operational costs per item processing

4.5 Compliance Assurance

Report	Grouping	Key Metrics & Focus
DLP Scrubbing Accuracy (Recurring)	Data Privacy	Biweekly sample accuracy vs. the Weeks 1–4 baseline; escalates to Simone Dubois / Legal on any drop below baseline
Works Council Logging Integrity	Employee Monitoring	Periodic audit confirming gate_decisions logs functional role only (e.g., Payment_PM), never individual employee IDs; escalates to Simone Dubois / Legal on any finding

5. 12-Week Monitoring Roadmap

Weeks 1–4 — Foundation

- **Baseline capture:** Capture all workflow/business and AI system baselines listed in Section 3.
- **Instrumentation:** set up DLP logging and complete the DLP latency audit; set the initial vector similarity threshold.
- **Shadow logging:** run the automated ingest in shadow mode alongside existing spreadsheets, without interrupting the team; compare AI-extracted fields (severity, product area, account tier) against what PMs enter manually.
- **Milestone:** $\geq 85\%$ alignment between AI tags and spreadsheet tags before Gate 1 goes live for real decisions.

Weeks 5–10 — Pilot Execution

1. **Ingestion & Gate 1:** PMs review incoming insights only via the Gate 1 interface. Track Gate 1 Acceptance vs. Override Rate. Refine feature-area taxonomy and calibrate similarity thresholds to reduce field-level overrides.

Milestones: ingestion lead time drops from ~20 minutes to < 1 minute per feedback item; RAG hit rate > 80%.

2. **Shortlisting & Gate 2:** automated scoring views replace spreadsheet-based prioritization. Gate 2 becomes a continuous per-item decision inside Appsmith.

Milestone: weekly triage meeting duration reduced by 75%.

3. **Drafting & Gate 3:** automated Jira ticket drafting and sync go live; Track Jira Draft Edit Rate and AI-vs-Reviewer story point variance.

Milestone: Gate 3 story acceptance rate > 85%.


Weeks 11–12 — Assessment

- **Spreadsheet retirement:** total PM time spent managing feedback drops > 80% vs. the Week 1–4 baseline (50% will also be acceptable considering the new SOP and UI).
- **Final measurement:** verify metrics against baseline; publish the full 12-week AI performance and cost/ROI analysis.
- **Decision input:** results measured against baseline feed the end-of-pilot stakeholder survey and the Proceed / Iterate / Stop decision (Brief §8).

6. Reporting Cadence & Ownership

Reporting is delivered continuously through the same dual architecture used elsewhere in the pilot: the Appsmith dashboard provides persistent, real-time views for standard executive monitoring (Product Analytics, Process & Workflow Analytics, AI System Analytics), while the /ask_data MCP agent supports ad-hoc, natural-language queries for any stakeholder who needs a specific cut of the data outside the standard reports.

Picture 1: Click on the diagram to open Feedback Intelligence & AI Analytics portal.



Ownership

Overall accountability for the Monitoring Plan sits with Elena Vance as Head of Product, while day-to-day operation and specialist oversight are distributed as follows:

Area	Owner	Responsibility
Overall Monitoring Plan	Elena Vance (Head of Product)	Accountable for the plan’s continued relevance, the Week 11–12 assessment, and the final Proceed / Iterate / Stop decision.
Day-to-day metric health & dashboard accuracy	Marcus Bauer (Senior PM)	Ensures data is flowing correctly, investigates anomalies, and keeps the operational reports usable for the team.

Area	Owner	Responsibility
AI system metrics (latency, cost, RAG quality, gate acceptance)	External Consultant + Priya Patel (Eng. Lead)	Maintains technical instrumentation and surfaces performance or cost issues.
Compliance & privacy metrics (DLP accuracy, role-only logging integrity)	Simone Dubois (Risk & Compliance)	Owens recurring DLP sample checks and Works Council logging audits; escalates any breach of baseline.
Baseline vs. actual comparison & final report	External Consultant + Elena Vance	Produces the Week 11–12 performance and ROI analysis that feeds the go/no-go decision.

7. Escalation Triggers

Automated BigQuery latency monitoring, with Appsmith/email notification triggers, watches gate turnaround against the 48-hour SLA escalation protocol (Risk Register OP-01), pending Elena Vance's executive agreement on that SLA.

A sustained drop below any Phase-level milestone in Section 5 (e.g., Gate 1 alignment below 85% past Week 4, or RAG hit rate below 80% past Week 6) is treated as a threshold-tuning trigger, prompting taxonomy or prompt adjustments before the next phase begins, rather than waiting for the Week 11–12 assessment.

8. End-of-Pilot Survey & Final Assessment

In Weeks 11–12, measured results are compared against the Weeks 1–4 baseline across both workflow/business and AI system metrics (Brief §8: “measured results against baseline; stakeholder Survey, Proceed / Iterate / Stop decision”).

The stakeholder confidence survey reuses the baseline instrument captured in Section 3 (“Does prioritization reflect real customer needs?”, 1–5) so the end-of-pilot score is directly comparable to the Week 1–4 baseline score.

Outputs feed the single Proceed / Iterate / Stop decision defined in the Project Brief, alongside the full 12-week AI performance and cost/ROI analysis.